

Music genre classification

March 19, 2022

In this project, the task is to classify the music genre of an audio track that is 30 seconds long. In the dataset, there is a total of ten music genres: *pop*, *metal*, *disco*, *blues*, *reggae*, *classical*, *rock*, *hip hop*, *country*, and *jazz*. The data is based on the **GTZAN dataset**¹, which consists of 1000 audio tracks, with 100 audio tracks from each music genre.




Figure 1: Wave plot of an audio track.

The model features, i.e., the input variables for the classifier, are commonly used audio features that are extracted from the audio track using the LIBROSA² package in Python (with default hyperparameter settings). Each feature has represented either its mean or standard deviation, indicated by the suffix in the feature name, e.g., *zero_cross_rate_mean*. The majority of the audio features are calculated from the spectrogram of the audio track, see Figure 2.

In the three data files *GenreClassData_5s.txt*, *GenreClassData_10s.txt*, and *GenreClassData_30s.txt*, respectively, you find features that have been calculated by splitting each audio track into non-overlapping segments of either 5, 10 or 30 seconds, before calculating the audio features, see Figure 3. If two rows have the same Track ID, they have been derived from the same audio track. In *Metadata_GenreClass.pdf* you find a detailed description of what the columns in the data matrix represent.

¹<http://marsyas.info/downloads/datasets.html>

²<https://librosa.org/doc/main/feature.html>




Figure 2: Spectrogram of an audio track.

The training set is represented by 800 audio tracks, with 80 audio tracks from each of the ten music genres. The test set is represented by 200 audio tracks, with 20 audio tracks from each of the ten music genres. The same audio tracks are placed in the same training and test set for all of the data files *GenreClassData_5s.txt*, *GenreClassData_10s.txt* and *GenreClassData_30s.txt*.




Figure 3: Audio track split into 5 second segments (green vertical lines) before calculating the features.

Task

1. For the first part use only the dataset *GenreClassData_30s.txt*.
 - (a) Design a k -NN classifier (with $k = 5$) for all ten genres using only the following four features: *spectral_rolloff_mean*, *mfcc_1_mean*, *spectral_centroid_mean* and *tempo*.
 - (b) Evaluate the performance of the classifier by finding the confusion matrix and the error rate for the test set.
 - (c) Discuss some of the misclassified tracks. Do you as a human agree with the classifier for some of the incorrect classifications?
2. Part 2 has focus on separability and its effect on the performance.
 - (a) For each of the four features; *spectral_rolloff_mean*, *mfcc_1_mean*, *spectral_centroid_mean* and *tempo*, compare the feature distribution for the four classes: **pop**, **disco**, **metal** and **classical**. Do this by producing histograms for each of the listed features and classes. Analyze how the feature distribution relates to the performance of your classifier by taking away the feature which shows most overlap between the classes. Train and test a classifier with the remaining three features.
 - (b) Compare the confusion matrices and the error rates for the four experiments and comment on the property of the features with respect to linear separability both as a whole and for the four separate classes.
3. Part 3 focuses on feature selection
 - (a) Design a k -NN classifier ($k=5$) for all ten genres using only four features with at least three features chosen among *spectral_rolloff_mean*, *mfcc_1_mean*, *spectral_centroid_mean*, and *tempo*.
 - (b) Motivate why you selected the particular four features.
 - (c) Generate the confusion matrix, compute the error rates and discuss the results.
4. Part 4 allows you to design your own classifier.
 - (a) Design a classifier for all ten genres that classifies the audio tracks, each represented by a *Track ID*. You are allowed to use any classifier, as many features as you like and all of the available data sets *GenreClassData_5s.txt*, *GenreClassData_10s.txt*, and *GenreClassData_30s.txt* as input data.
 - (b) Generate the confusion matrix, compute the error rates
 - (c) Justify your choices and discuss the results.